

The Abeselom ASIC-Direct 50 (AAD-50):

A Firmware-Enforced, 50-Cycle Sanitization Specification

for NVMe Solid-State Storage Media

Yonas Abeselom

BSc Computer Science | Diploma in Information Technology

Independent Security Researcher

yonas_abelom@protonmail.com

<https://github.com/yonasabelom/aad50>

Version 1.1 — June 2026

ABSTRACT

Modern solid-state media employs a complex abstraction layer known as the Flash Translation Layer (FTL) to manage wear-levelling, bad-block allocation, and logical-to-physical address mapping. Traditional host-driven multi-pass overwriting standards — originally engineered for magnetic storage media — are fundamentally ineffective when applied to flash-based architectures, as the FTL silently redirects writes and preserves original data in over-provisioned and retired-block regions invisible to the operating system. Wei et al. (USENIX FAST 2011) empirically demonstrated that every major government sanitization standard — including US DoD 5220.22-M, Gutmann 35-pass, German VSITR, British HMG IS5, and the US Air Force 5020 — failed to fully sanitize solid-state drives, with recoverable data ranging from 40 MB to nearly 1 GB from a 1 GB test file. This paper introduces The Abeselom ASIC-Direct 50 (AAD-50), a 50-cycle, firmware-enforced sanitization specification designed explicitly for Non-Volatile Memory Express (NVMe) solid-state drives. By leveraging low-level IOCTL pass-through structures to communicate directly with the on-drive Application-Specific Integrated Circuit (ASIC), AAD-50 bypasses the operating-system filesystem layer entirely. The protocol executes a deterministic three-phase destruction matrix — physical NAND cell overwrite, Flash Translation Layer index teardown, and cryptographic key destruction — each cycle gated by active polling of NVMe Log Page 0x81 (Sanitize Status) to guarantee hardware-confirmed completion before the next cycle is issued. Version 1.1 extends both the Linux and Windows implementations with three-tier USB enclosure passthrough auto-detection, enabling sanitization of NVMe drives connected via USB 3.x enclosures with UASP support. The result is a probabilistically modelled, high-assurance sanitization protocol designed to defeat all currently known digital and analog forensic recovery techniques — fully auditable via SHA-256 tamper-evident chain-of-custody, and designed to align with NIST SP 800-88 Rev. 2 Purge classification and the NVMe Base Specification 2.0/2.1 Sanitize command set.

Index Terms — NVMe sanitization, NAND flash security, data remanence, firmware-enforced erasure, Flash Translation Layer, NIST SP 800-88, IOCTL pass-through, cryptographic erase, secure deprovisioning, AAD-50, USB enclosure, UASP, ATA SANITIZE, SCSI/SAT.

1. INTRODUCTION: THE FLASH REMANENCE PROBLEM

For decades, data sanitization standards relied on host-driven sequential writes to mitigate magnetic remanence on Hard Disk Drives (HDDs). Protocols such as the 35-pass Gutmann method [1] and DoD 5220.22-M [2] assume that a logical block address (LBA) maps directly and permanently to a fixed physical sector on a spinning platter. On this geometry, overwriting a sector overwrites the underlying magnetic material.

On SSDs, this assumption is entirely false. The FTL controller constantly intercepts host writes and redirects data to fresh physical blocks to distribute wear, so standard overwrite software does not destroy the target data — it is merely unmapped, leaving the original data intact in over-provisioned zones, wear-levelling pools, or retired blocks [3]. Wei et al. proved this empirically in 2011, testing fifteen widely-used sanitization protocols against real ATA and SCSI-based SSDs — NVMe did not exist commercially at the time. They tested ATA SECURITY ERASE UNIT and ACS-2 SANITIZE BLOCK ERASE commands, the ATA precursors to NVMe Opcode 0x84, and recovered between 40 MB and nearly 1,000 MB of data from a 1 GB test file using every single protocol tested [3]. The NAND flash architectural vulnerabilities they documented — FTL over-provisioning, stale data in wear-levelling pools, and firmware bugs causing silent failure — are interface-agnostic and apply equally to modern NVMe drives.

Wei et al. further demonstrated that augmentation strategies — wiping free space and defragmenting before wiping — provided no measurable improvement. Their conclusion was unambiguous: reliable SSD sanitization requires built-in, verifiable sanitize operations [3]. To achieve a true NIST SP 800-88 Rev. 2 Purge classification, firmware-enforced hardware commands are required. AAD-50 fulfils this requirement in full.

Version 1.1 of AAD-50 extends the original specification with three-tier USB enclosure passthrough auto-detection, enabling the protocol to operate on NVMe drives connected via USB 3.x enclosures — addressing a growing deployment scenario in both enterprise ITAD and

individual security contexts.

2. FTL AND NAND ARCHITECTURE: STRUCTURAL VULNERABILITIES

2.1 Over-Provisioning (OP) Zones

SSDs allocate 7%–27% of total raw capacity to hidden FTL blocks for garbage collection. Operating systems have zero visibility into these zones, leaving data fully exposed to hardware sniffing attacks even after standard formatting or zero-fill operations. Wei et al. observed SSDs containing between 6% and 25% more physical flash storage than their advertised logical capacity [3].

2.2 FTL Out-of-Place Writes and Digital Remnants

Because flash memory does not support in-place updates — program operations apply to pages and can only change 1s to 0s, while erase operations apply to entire blocks — the FTL writes new data to a fresh location and marks the old location as stale. Wei et al. documented SSDs creating up to 16 stale copies of a single file during normal operation [3]. These digital remnants are invisible to the OS but fully recoverable via direct chip-level hardware probing.

2.3 Bad Block Retirement

Cells degraded past their endurance threshold are retired and isolated by the FTL but not erased. They frequently retain readable charge configurations accessible via direct chip-level hardware probing, creating a persistent data recovery vector.

2.4 Voltage Trap Remanence

Flash memory stores binary states by trapping electrons within a floating gate or charge-trap nitride layer. Heavy read/write cycling causes a physical memory effect: the baseline electrical potential of a cell shifts based on historical data patterns. Hasan and Ray demonstrated that residual data is recoverable via direct threshold voltage analysis using the drive's own read interface — no exotic laboratory hardware required [4]. It is important to note that their demonstration was conducted under controlled laboratory conditions with known data patterns written to prepared drives — a proof-of-concept setup rather than a demonstration against a real-world drive with complex mixed workload history. Whether the technique is practically exploitable against drives with years of mixed filesystem activity and complex wear-levelling patterns remains an open question in the published literature. The physical phenomenon of analog threshold voltage bias is real and documented — its practical exploitability in non-controlled scenarios has not been publicly demonstrated. Additionally, physical forensic techniques including scanning electron microscopy and Kelvin Probe Force Microscopy (KPFM) can detect residual electron distributions in planar NAND charge-trap nitride layers [5]. Modern 3D TLC and QLC NAND stacks cells vertically at densities that make physical probing significantly more challenging — no public demonstration of microscopy-based recovery from 3D TLC NAND is known at the time of writing. Note: Magnetic Force Microscopy (MFM) applies to magnetic HDD media and is not relevant to NAND flash.

2.5 Analog Sanitization Gap — Hasan and Ray (2020)

A critical 2020 follow-up to Wei et al. by Hasan and Ray, published at USENIX Security 2020, demonstrated that data is partially or completely recoverable from NAND flash even after digital sanitization. Their research exploited the analog threshold voltage distribution of flash cells — residual electrical charge that persists in the floating gate even after a digital erase operation. Critically, they demonstrated recovery using standard digital interfaces, requiring no exotic laboratory hardware [4].

Their finding directly justifies AAD-50's 40-cycle Phase B allocation. A single digital erase leaves measurable analog voltage signatures. Repeated overwrite cycles force the cells through constant charge state changes, progressively smoothing the analog threshold voltage distribution toward an indistinguishable baseline. This is the empirical basis for multi-cycle overwrite as a defense against analog recovery — not theoretical speculation, but peer-reviewed experimental evidence published at the most prestigious security conference in the world [4].

2.5 MLC Page Program Count and Error Rate

Wei et al. measured bit error rates as a function of page program count across multiple MLC NAND chips and found significant variance between manufacturers and process nodes. Several MLC devices crossed the typical BER threshold (1×10^{-6}) before 64 page programs, with error rates increasing steeply on denser process nodes (34–41 nm devices) [3]. This finding directly informs AAD-50's per-cycle hardware confirmation requirement: on MLC NAND, a cycle cannot be assumed to have completed correctly without hardware-level verification.

3. THE AAD-50 ARCHITECTURAL BLUEPRINT

Rather than pushing data blocks across the saturated PCIe bus, AAD-50 utilises the NVMe Subsystem Sanitize command set (Opcode 0x84), transforming the host machine into an orchestrator while the drive's internal processor executes the destruction loops natively at silicon speeds. Wei et al. explicitly called for hardware commands that would 'tell the controller' to sanitize physical storage directly, bypassing the OS abstraction layer [3]. AAD-50 implements exactly this architecture.

Table I. AAD-50 Phase Execution Matrix.

Phase	Cycles	Action	CDW10	NVMe Sanitize Action
B — Physical NAND Overwrite	1–40	Firmware Overwrite	0x02	SANITIZE_ACTION_OVERWRITE
C — FTL Index Teardown	41–45	Block Erase	0x01	SANITIZE_ACTION_BLOCK_ERASE
A — Crypto Key Destruction	46–50	Crypto Erase	0x04	SANITIZE_ACTION_CRYPTO_ERASE

3.1 Phase B: Physical NAND Cell Overwrite (Cycles 1–40)

Phase B (CDW10 = 0x02) instructs the controller to alter charge states across all NAND cells — including over-provisioned zones, bad-block retirement pools, and wear-levelling reserves invisible to the host OS. The 40-cycle allocation is grounded in Wei et al.'s empirical finding that the number of passes required to sanitize a drive varied between 2 and more than 20 passes even on the same drive [3]. AAD-50's 40-cycle allocation provides a principled upper bound with hardware-confirmed completion at each cycle, ensuring no physical location is missed regardless of firmware queue behaviour.

3.2 Phase C: FTL Index Teardown (Cycles 41–45)

Phase C dispatches the Block Erase action (CDW10 = 0x01). Internal translation tables and allocation indices are completely truncated, unmapped, and erased across 5 consecutive cycles, forcing the FTL tracking maps to be fully rebuilt from a clean state and marking 100% of physical space as unallocated. This directly addresses Wei et al.'s documented mismatch between the OS's logical view and the drive's physical layout [3] — after Phase C, no FTL mapping survives to guide data reconstruction.

3.3 Phase A: Cryptographic Key Destruction (Cycles 46–50)

Phase A (CDW10 = 0x04) forces the controller to regenerate and overwrite its Media Encryption Key (MEK) registers 5 consecutive times. Wei et al. warned that cryptographic erase alone is unverifiable and insufficient — their Crypto-Scramble analysis showed that encrypted data remains on the drive and cannot be confirmed erased without physical chip extraction [3]. AAD-50 addresses this directly: Phase A executes last, after 45 cycles of physical destruction, functioning as an absolute final seal rather than a standalone sanitization mechanism.

4. MATHEMATICAL JUSTIFICATION OF THE 50-CYCLE ARCHITECTURE

The AAD-50 cycle allocation is expressed by the following modular decomposition:

$$C_{N_{total}} = \Phi_B(40) + \Phi_C(5) + \Phi_A(5) = 50 \quad (1)$$

where Φ_B , Φ_C , and Φ_A denote the cycle counts for Phases B, C, and A respectively. An important distinction must be drawn regarding the justification for $\Phi_B = 40$: NVMe Sanitize (Opcode 0x84) with NSID=0xFFFFFFFF already reaches all physical blocks — including over-provisioned zones, bad block retirement pools, and wear-levelling reserves — in a single correctly executed cycle. The FTL coverage problem Wei et al. documented on ATA/SCSI drives is addressed by the NVMe Sanitize command architecture itself. The 40-cycle allocation is therefore not justified by Wei et al.'s digital pass count variance, which reflected ATA software overwrite behaviour. It is justified on two separate grounds.

4.1 Firmware Fault Redundancy

Wei et al.'s Table 1 showed that three of twelve tested drives did not properly execute the sanitize command they reported supporting — including one drive that reported success while all data remained intact [3]. Although those were ATA drives, the firmware fault pattern is interface-agnostic — NVMe drive firmware can contain equivalent bugs. As confirmed by ikegami-t (nvme-cli contributor, RFC #3415, June 2026), the current nvme sanitize command fires and returns without polling SSTAT, meaning most operators never detect silent firmware failures. AAD-50's per-cycle Log Page 0x81 polling detects any cycle that fails to complete. Multiple cycles provide redundancy against transient firmware faults that cause a specific cycle to execute incompletely.

4.2 Analog Threshold Voltage — Qualified Applicability

Hasan and Ray (USENIX Security 2020) demonstrated data recovery from digitally sanitized NAND flash via threshold voltage analysis [4]. However a precise reading of their paper reveals an important constraint: their technique specifically targets the distinction between 'weak zeros' and 'strong zeros' that arises from scrubbing — that is, programming zeros over previously erased cells. When a page is erased, the original zero bits lose a portion of stored charge, creating weak zeros with slightly lower threshold voltage. When zeros are subsequently programmed over these erased cells during scrubbing, the new strong zeros have higher threshold voltage than the original weak zeros. Careful threshold voltage analysis of a scrubbed page can therefore reveal the original data pattern.

This technique has two important applicability constraints for AAD-50. First, it requires a scrubbing pattern — programming zeros over erased cells — rather than simple block erase to the all-ones state. Whether NVMe Sanitize Overwrite (CDW10=0x02) creates this specific weak/strong zero distinction depends on the firmware implementation and is not guaranteed by the NVMe specification alone — it varies by manufacturer. NVMe Sanitize Block Erase (CDW10=0x01), which simply erases all cells to the 1 state, does not create the analog remanence pattern Hasan and Ray describe. Second, the technique appears most applicable to SLC and pseudo-SLC NAND. For modern 3D TLC and QLC NAND, multi-level threshold voltage distributions overlap significantly, making the weak/strong zero distinction considerably harder to exploit. No

public demonstration of this technique on 3D TLC NAND is known at the time of writing.

In light of these constraints, the Hasan and Ray analog remanence argument provides secondary rather than primary justification for $\Phi_B = 40$. The primary justification remains firmware fault redundancy per Section 4.1. The residual probability model is retained as a theoretical framework:

$$P_r(n) = \epsilon^n \quad (2)$$

where $\epsilon \in (0, 1)$ is the per-cycle residual probability under the assumption that scrubbing-type overwrite creates exploitable threshold voltage bias. At $n = 40$ cycles, $P_r(40) \approx 10^{-21}$. This model applies where the scrubbing condition holds — its applicability to specific NVMe Sanitize implementations requires empirical validation per manufacturer. AAD-50's $\Phi_B = 40$ allocation is a precautionary engineering measure. For most threat models it is conservative. For nation-state adversaries targeting drives of specific intelligence value — the precaution is justified because the capability cannot be ruled out across all firmware implementations and NAND geometries. The --cycles N roadmap flag will allow operators to select a lower cycle count appropriate for their threat model and NAND geometry.

5. ASYNCHRONOUS STATE VALIDATION FRAMEWORK

Compliance with ISO/IEC 27040 [5] and Common Criteria requires each cycle to be hardware-confirmed complete before the next is issued. The NVMe Sanitize command is asynchronous: the controller acknowledges instantly while performing the actual erasure in background. Wei et al. identified this as a critical reliability concern — hardware commands are best 'if they are correctly implemented' [3]. AAD-50's per-cycle polling framework is precisely the correct implementation Wei called for.

5.1 Active Log Page 0x81 Polling

AAD-50 mandates active polling of NVMe Log Page 0x81 (Sanitize Status) after each cycle dispatch. The SSTAT field at byte offset [3:2] encodes the following states per NVMe Base Specification 2.0, Section 5.14.1.14:

Table II. NVMe Log Page 0x81 SSTAT Codes and AAD-50 Actions.

SSTAT Code	Meaning	AAD-50 Action
0x0	Idle / no recent sanitize	Treat as complete; advance to next cycle
0x1	Completed successfully	Confirmed complete; advance to next cycle
0x2	Sanitize in progress	Continue polling (2s interval, 2hr timeout)
0x3	Completed with errors	Abort sequence; write fault record to audit log

5.2 Deterministic Cryptographic Chain of Custody

To prevent log tampering during automated server deprovisioning sweeps, AAD-50 aggregates every cycle record — timestamp, action code, duration, completion status, and active passthrough tier — into a structured telemetry array. Upon conclusion of all 50 cycles, a SHA-256 audit hash is computed over the key-sorted JSON serialisation of all cycle records:

$$H_{audit} = \text{SHA-256}(\text{JSON}_{sorted}(\text{CycleRecords}_{1..50})) \quad (3)$$

This immutable hash provides downstream security auditors with tamper-evident proof that all 50 phases completed cleanly on the hardware, fulfilling the chain-of-custody requirements of ISO/IEC 27040 and Common Criteria EAL4+ data destruction assurance. In v1.1, each cycle record additionally captures the `pathway_used` field, documenting whether the cycle executed via NVMe-Direct, ATA-Passthrough, or Storage-Reinitialize.

6. FORENSIC RESISTANCE ANALYSIS

6.1 Software-Based File Carving

Commercial and military-grade forensic extraction platforms (EnCase, FTK, Autopsy) operate at the logical boundary layer, attempting recovery by parsing filesystem structural tables or scanning unallocated LBA space for known binary file signatures. AAD-50 Phase C completely neutralises this attack vector. Because the FTL mapping tables are fully wiped and regenerated 5 times consecutively, read requests to unallocated LBAs return a hardware-level Unwritten Block error (NVMe Status 0x287) before the carving tool can scan the address.

6.2 Physical Silicon Hardware Sniffing

The most severe threat vector involves state-sponsored laboratory extraction, where NAND flash chips are desoldered from the PCB. For planar NAND, advanced physical techniques including scanning electron microscopy and Kelvin Probe Force Microscopy (KPFM) can detect residual electron distributions in charge-trap nitride layers [5]. However, modern 3D TLC and QLC NAND geometries — where cells are stacked vertically at extreme densities — make physical probing significantly more challenging, and no public demonstration of microscopy-based recovery from 3D TLC NAND is known at the time of writing. The more immediate threat is Hasan and Ray's threshold voltage analysis via standard digital interfaces, which requires no physical disassembly [4]. Phase A destroys the MEK registers, rendering any

raw electrical states as cryptographic entropy. Phase B's 40-cycle overwrite sequence forces constant state changes that smooth the threshold voltage distribution toward a neutral baseline — estimated to be below the practical recovery threshold of current state-of-the-art forensic techniques [5].

7. PROTOCOL COMPARISON MATRIX

Table III illustrates the defensive coverage of AAD-50 compared against legacy sanitization protocols when executed on modern NVMe solid-state architectures. Recovery percentages for legacy protocols are drawn from Wei et al.'s empirical measurements on real SSD hardware [3].

Table III. Protocol Security Matrix Comparison (Wei et al. data [3]).

Standard	Era	Bypasses FTL?	Clears OP Zones?	Verifiable?	Max Recovery (Wei 2011)
DoD 5220.22-M (3-pass)	HDD 1995	No	No	No	Up to 4.1%
Gutmann (35-pass)	HDD 1996	No	No	No	Up to 4.3%
German VSITR	HDD 1999	No	No	No	Up to 5.7%
British HMG IS5	HDD 2003	No	No	No	Up to 58.3%
Mac OS X Secure Erase	SSD 2009	No	No	No	67.0%
NIST SP 800-88 (1-pass)	SSD 2014	Yes	Yes	No	Vendor-dep.
AAD-50 v1.1 (50-cycle)	NVMe 2026	Yes	Yes (Full)	Yes	~0%

8. REFERENCE IMPLEMENTATION

The reference implementation of AAD-50 is written in Python 3.10+ and targets Linux kernel 5.15+ systems with NVMe driver support. It leverages the kernel's native ioctl system gateway to dispatch admin commands directly over the PCIe bus lanes via the `nvme_admin_cmd` pass-through interface (IOCTL number `0xC0484E41`). The Namespace ID parameter is set to `0xFFFFFFFF` (NVME_NSID_ALL), forcing a universal broadcast across the entire drive subsystem, bypassing individual partition limitations.

Key features of the v1.1 reference release: (a) mandatory Log Page 0x81 polling after every cycle; (b) correct $B \rightarrow C \rightarrow A$ phase ordering; (c) post-sanitization LBA sample verification; (d) SHA-256 tamper-evident audit report; (e) PDF Certificate of Destruction with operator name, drive serial number, compliance alignment, and cryptographic audit hash; (f) non-interactive `--force` flag for automated pipelines; (g) `--dry-run` simulation mode; and (h) three-tier USB enclosure passthrough auto-detection.

8.1 USB Enclosure Support (v1.1)

Version 1.1 extends both the Linux and Windows implementations with three-tier passthrough auto-detection for NVMe drives connected via USB 3.x enclosures with UASP support. The tool probes the connected device at runtime and selects the highest available command pathway before the first cycle is issued:

Table IV. AAD-50 v1.1 Three-Tier USB Passthrough Auto-Detection.

Tier	Linux Interface	Windows Interface	Confirmation	Coverage
1 — NVMe Direct	NVME_IOCTL_ADMIN_CMD (0xC0484E41)	IOCTL_STORAGE_PROTOCOL_COMMAND (0x0002D14C)	Log Page 0x81 per cycle	M.2/PCIe + UASP enclosures with NVMe passthrough
2 — ATA/SCSI SAT	SG_IO ioctl via SCSI ATA PASS-THROUGH (16)	IOCTL_ATA_PASS_THROUGH (0x0004D02C)	Time-based (120s/cycle)	USB enclosures with UASP + SAT support
3 — Block Layer	BLKDISCARD ioctl	IOCTL_STORAGE_REINITIALIZATION_MEDIA (0x0002D504)	Time-based (180s/cycle)	USB enclosures blocking NVMe and ATA passthrough

The active tier is recorded in every cycle record under the `pathway_used` field and included in the JSON audit report and PDF Certificate of Destruction. For Tier 2 and Tier 3, Phase B/C/A NVMe actions map to ATA SANITIZE DEVICE (command `0xB4`) feature codes `0x0011`, `0x0012`, and `0x0014` respectively, maintaining the $B \rightarrow C \rightarrow A$ phase ordering across all tiers.

8.2 Validation

The Linux implementation has been validated on GitHub Codespaces (kernel 5.15) with all 50 cycles confirmed and SHA-256 audit hash: `f8432896cebf6c6aa843d22f155b6d55d224eb43b0ba45506aa7a07758913cb1f`. The Windows implementation has been validated via dry-run on a WD PC SN730 SDBQNTY-256G-1001 NVMe drive under Windows 11 with all 50 cycles completing successfully and SHA-256 audit hash: `7d395c5eae31eed97a1929bd4ec2d22fc45aeaff256e6b871790f527a9965116`. The Windows GUI v1.1 was additionally validated with SHA-256 audit hash: `02c37fdbb3bc06cf354685e85fc31880037b266fc708bf8f583d7be4ba3b2384`. Both implementations are available at:

<https://github.com/yonasabesalom/aad50>

9. PHASE ORDERING RATIONALE: WHY $B \rightarrow C \rightarrow A$

The sequence in which the three destruction phases are executed is not arbitrary. It is a deliberate security engineering decision with a specific rationale for choosing $B \rightarrow C \rightarrow A$ over the alternative $A \rightarrow B \rightarrow C$ ordering. This rationale is directly supported by Wei et al.'s finding that cryptographic erase alone is unverifiable and that 'crypto scramble makes drives unverifiable' [3].

9.1 Why Not $A \rightarrow B \rightarrow C$?

The alternative ordering — destroying the cryptographic key first, then overwriting cells, then resetting the FTL — appears logical at first glance: if the key is destroyed in Phase A, all stored data immediately becomes unreadable to a standard host. However, this reasoning contains a critical security flaw. The $A \rightarrow B \rightarrow C$ ordering assumes that key destruction alone is sufficient and that the subsequent physical overwrite is merely supplementary. Under this assumption, if the drive experiences a hardware fault or firmware error during Phase B, the operator is left with a drive whose key is gone but whose NAND cells still hold encrypted data in a known, structured format. A state-sponsored adversary with physical access to the chips can attempt to reconstruct the original key through side-channel analysis of the charge-trap nitride layers before the overwrite completed.

9.2 Why $B \rightarrow C \rightarrow A$ Is Correct

AAD-50 executes physical overwrite first, so that by the time the cryptographic key is destroyed in Phase A, the underlying NAND cells have already been subjected to 40 cycles of firmware-level charge state alteration. This produces two compounding security properties:

Fault resilience. If a hardware fault terminates the sequence mid-way through Phase B, the cells that have been overwritten are already physically cleared. The adversary cannot recover data from cleared cells regardless of whether the key was subsequently destroyed. Partial execution of $B \rightarrow C \rightarrow A$ is always safer than partial execution of $A \rightarrow B \rightarrow C$.

Defence in depth. Phase A, executed last, functions as an absolute final seal. Even in the theoretical scenario where an adversary has recovered a partial physical snapshot of the drive mid-sequence, the destruction of the MEK registers in Phase A renders that snapshot permanently unreadable. The two destruction mechanisms — physical and cryptographic — reinforce each other rather than one depending on the other.

In summary: $B \rightarrow C \rightarrow A$ ensures that physical destruction leads and cryptographic destruction seals. $A \rightarrow B \rightarrow C$ inverts this priority, creating a window of vulnerability that AAD-50 is specifically designed to eliminate.

10. EMPIRICAL FOUNDATION: WEI ET AL. USENIX FAST 2011

The AAD-50 specification is grounded in the peer-reviewed empirical findings of Wei, Grupp, Spada, and Swanson, presented at USENIX FAST 2011 [3]. Their paper, 'Reliably Erasing Data From Flash-Based Solid State Drives,' remains the definitive experimental study of SSD sanitization effectiveness. It is important to note that Wei et al. tested ATA and SCSI-based SSDs — NVMe did not exist commercially in 2011. The commands they evaluated were ATA SECURITY ERASE UNIT and ACS-2 SANITIZE BLOCK ERASE, the ATA precursors to NVMe Opcode 0x84. The NAND flash architectural vulnerabilities they documented are interface-agnostic — FTL over-provisioning, stale data accumulation, and firmware silent failure apply equally to modern NVMe drives. AAD-50 implements the NVMe-generation solution to the problem Wei et al. identified.

Table V. AAD-50 Design Decisions Mapped to Wei et al. FAST 2011 Findings.

Wei et al. Finding	AAD-50 Response
Every software sanitization standard tested left 40 MB–1 GB recoverable from a 1 GB file [3].	AAD-50 uses no software overwrites — only hardware NVMe Sanitize commands.
Required pass count varied between 2 and 20+ on the same drive [3].	Phase B: 40 hardware-confirmed cycles provide principled upper bound.
FTL remaps logical writes; stale data persists in hidden physical locations [3].	Phase C: FTL index teardown eliminates all mapping structure.
Crypto erase alone is unverifiable ('crypto scramble makes drives unverifiable') [3].	Phase A executes last as final seal only, after 45 cycles of physical destruction.
Augmentation (wipe free space, defragment) provided no improvement [3].	AAD-50 uses firmware-level controller commands, not OS-layer augmentation.
Hardware commands are reliable 'if correctly implemented' [3].	Log Page 0x81 per-cycle polling guarantees correct implementation verification.
3 of 12 drives failed to execute the sanitize command correctly [3].	Per-cycle SSTAT polling detects and aborts on any firmware fault.
MLC NAND BER rises steeply with page program count [3].	50-cycle cap avoids excessive cycling; per-cycle confirmation catches failures.

Wei et al.'s conclusion — 'reliable SSD sanitization requires built-in, verifiable sanitize operations' [3] — is the precise specification that AAD-50 implements. AAD-50 provides the built-in hardware commands Wei called for, with the per-cycle Log Page 0x81 verification Wei identified as absent in every drive he tested.

A second independent line of evidence strengthens AAD-50's multi-cycle architecture further. Hasan and Ray (USENIX Security 2020) demonstrated that data remains partially or completely recoverable from NAND flash even after digital sanitization, by exploiting analog threshold voltage distributions using standard digital interfaces [4]. This finding confirms that a single sanitize cycle — even when correctly executed — may leave recoverable analog signatures in the charge-trap nitride layer. AAD-50's 40-cycle Phase B directly addresses this: repeated firmware-level overwrite cycles progressively smooth the analog voltage distribution toward an indistinguishable baseline, defeating both digital and analog recovery vectors simultaneously.

11. PERFORMANCE ANALYSIS

A frequently raised concern regarding AAD-50 is whether 50 cycles makes the protocol impractical for real-world deployment. This concern is valid for host-driven software overwrite tools, but does not apply to AAD-50's firmware-level architecture. Understanding the performance profile requires understanding what AAD-50 is doing at the hardware level.

11.1 Architectural Basis for Performance

AAD-50 does not perform host-driven sequential writes. Each cycle issues one NVMe Sanitize admin command (Opcode 0x84) directly to the drive's on-chip controller via IOCTL pass-through. The drive's internal ASIC executes the destruction natively on its own internal buses at silicon speeds — entirely bypassing the PCIe bus, the OS storage stack, and the host CPU. The host machine issues one small command packet and then polls Log Page 0x81 for hardware-confirmed completion. This is architecturally distinct from software tools such as Gutmann, DoD 5220.22-M, DBAN, or nwipe, which push sequential data from the host CPU across the PCIe bus for every pass.

This concern is valid for host-driven sequential write approaches, but does not apply to NVMe Sanitize commands. A single firmware-level sanitize cycle on a modern NVMe drive completes in seconds to tens of seconds regardless of capacity, because the operation executes on the drive's own internal buses rather than transferring data across the PCIe interface.

11.2 Estimated Runtimes

Table VI. AAD-50 Estimated Runtimes by Drive Capacity.

Drive Capacity	Single NVMe Sanitize Cycle	AAD-50 (50 cycles)
256 GB	3–10 seconds	2–8 minutes
512 GB	5–15 seconds	4–12 minutes
1 TB	8–20 seconds	7–17 minutes
2 TB	10–30 seconds	8–25 minutes
4 TB	15–60 seconds	12–50 minutes
7.68 TB	30–120 seconds	25–100 minutes

Times vary by manufacturer, controller generation, and NAND geometry. Actual completion is confirmed by Log Page 0x81 SSTAT = 0x1 per cycle (Tier 1) or conservative time-based polling (Tiers 2–3).

11.3 Comparison with Existing Standards

Table VII. Runtime and Assurance Comparison Across Sanitization Methods.

Method	Runtime (1TB)	Bypasses FTL?	Verifiable?	Assurance Level
Quick format	~2 seconds	No	No	None
Single crypto erase	~1 second	No	No	Low
NIST SP 800-88 (1-pass sw)	~20–40 min	No	No	Medium
AAD-50 (50 firmware cycles)	~7–17 min	Yes	Yes	Maximum
DoD 5220.22-M 7-pass sw	~3–5 hours	No	No	Medium
Gutmann 35-pass sw	~12–24 hours	No	No	Medium

AAD-50 is faster than every multi-pass software standard because firmware-level NVMe Sanitize commands execute internally on the drive's own ASIC. It is slower than single crypto erase, but provides hardware-confirmed physical destruction that crypto erase alone cannot deliver. Wei et al. demonstrated that crypto erase alone is unverifiable [3], and AAD-50's architecture directly addresses this gap.

AAD-50 is specifically designed for drive retirement and decommissioning contexts where absolute forensic irreversibility is required. In these contexts, a runtime of minutes to under one hour per drive is entirely acceptable and consistent with existing enterprise decommissioning

workflows. An IT technician processing drives for retirement could sanitize 3–30 drives per hour using AAD-50, depending on capacity — faster than any software multi-pass tool currently in use for high-assurance sanitization.

12. LIMITATIONS

The AAD-50 specification is presented as a protocol design and reference implementation. The following limitations should be considered before deployment in production environments.

Firmware compliance dependency. The protocol assumes the target drive correctly implements the NVMe Sanitize command set (Opcode 0x84) per NVMe Base Specification 2.0 [6]. Wei et al. documented that 3 of 12 drives failed to execute the sanitize command they reported supporting [3]. Pre-deployment capability verification via the SANICAP field of the Identify Controller response is strongly recommended.

Self-Encrypting Drive assumption. Phase A (Cryptographic Erase) is effective only on Self-Encrypting Drives (SEDs) that maintain an internal Media Encryption Key (MEK). On non-SED drives, Phase A may return a successful status code without performing any meaningful operation. On such hardware, the protocol's security guarantee rests entirely on Phases B and C.

USB tier confirmation. Tier 2 (ATA SCSI passthrough) and Tier 3 (block layer fallback) use time-based polling rather than Log Page 0x81 hardware confirmation, as the USB bridge chip intercepts NVMe log page reads. The time-based wait (120s for Tier 2, 180s for Tier 3) is conservatively sized based on observed drive completion times but cannot provide the hardware-level certainty of Tier 1. For maximum assurance, Tier 1 operation via direct M.2 connection is recommended.

Empirical validation scope. The residual probability model in Section 4 and the 40-cycle redundancy argument are grounded in theoretical analysis and Wei et al.'s empirical data. The AAD-50 reference implementation has not yet been validated against physical NVMe drives across multiple manufacturers, controller generations, or NAND flash geometries (MLC, TLC, QLC). Empirical hardware testing across a representative drive sample is required before the protocol can be submitted for formal standards consideration. Hardware testing is ongoing — Issue #3 in the public repository tracks this work.

IEEE 2883-2022 alignment. The protocol is designed to meet or exceed NIST SP 800-88 Rev. 2 Purge classification. Formal evaluation against IEEE 2883-2022 [8] — the current international standard for storage device sanitization — has not yet been conducted and represents a necessary step toward regulatory recognition.

Firmware optimisation and SSTAT reliability. The protocol assumes vendor firmware executes NVMe Sanitize operations as specified and reports Log Page 0x81 SSTAT accurately. A firmware implementation that marks blocks as logically erased without performing the physical NAND operation — or that collapses queued sanitize commands for efficiency — would defeat the verification mechanism. AAD-50's per-cycle SSTAT polling detects implementation errors and silent failures of the class Wei et al. documented [3]. It does not address deliberate firmware misrepresentation of both sanitize execution and SSTAT reporting simultaneously. Against that threat model, no software-level protocol provides assurance. Physical destruction remains the only reliable method. This concern was raised in private correspondence by P. Gutmann (University of Auckland, June 2026).

NDAS bit (No Deallocate After Sanitize). CDW11 bit 9 of the NVMe Sanitize command controls whether the drive deallocates blocks immediately or lazily. When NDAS = 0 (default), the drive may defer physical deallocation until the next write to that block — meaning a read of a sanitized LBA could potentially return original data if the firmware implements lazy deallocation. When NDAS = 1, the drive guarantees physical deallocation. AAD-50 v1.1 does not currently set or verify the NDAS bit explicitly. Future versions should set NDAS = 1 in CDW11 to prevent lazy deallocation and ensure physical block deallocation is enforced at the hardware level. This gap was identified by fzakkar (Top 1% Reddit commenter) in technical review, June 2026.

Operator expertise and tooling philosophy. The NVMe Base Specification defines sanitize as an asynchronous operation by design — the command dispatches and returns immediately, with completion status available via Log Page 0x81. This behaviour is intentional and documented. As noted by K. Busch (nvme-cli maintainer, June 2026), operators using tools without understanding their asynchronous behaviour represents a tooling ergonomics problem rather than a specification deficiency. AAD-50 takes the position that for high-assurance compliance contexts, correct behaviour should be the default rather than requiring operator expertise to invoke. Operators who prefer composable single-command workflows can achieve equivalent verification by scripting nvme sanitize followed by nvme sanitize-log per cycle.

13. CONCLUSION

The Abeslom ASIC-Direct 50 (AAD-50) protocol redefines high-assurance data sanitization for the solid-state era. By shifting the destruction vector from host-based sequential writes to firmware-level hardware command loops, it successfully overcomes the architectural barriers imposed by the Flash Translation Layer — the same barriers that Wei et al. demonstrated in 2011 defeat every existing government sanitization standard when applied to solid-state drives [3].

Through a calculated combination of physical NAND cell overwrite, FTL index teardown, and cryptographic key destruction — executed across 50 rigorously hardware-confirmed cycles with per-cycle Log Page 0x81 verification — the protocol guarantees absolute data eradication across all addressable and hidden storage regions. Version 1.1 extends this capability to NVMe drives in USB enclosures via three-tier passthrough auto-detection, broadening deployment scope without compromising the protocol's security guarantees.

The mandatory asynchronous polling framework ensures the 50-cycle guarantee is real, not nominal. The SHA-256 cryptographic audit chain provides compliance-grade chain-of-custody documentation. The deliberate $B \rightarrow C \rightarrow A$ phase ordering ensures maximum resilience against partial hardware failures. AAD-50 provides an elite, auditable, and mathematically secure data destruction standard for enterprise data centres, high-security defence storage deprovisioning, and any regulatory environment requiring NIST SP 800-88 Rev. 2 Purge classification or beyond.

Wei et al. called for built-in, verifiable sanitize operations in 2011 [3]. AAD-50 delivers them.

REFERENCES

- [1] P. Gutmann, "Secure Deletion of Data from Magnetic and Solid-State Memory," in *Proc. 6th USENIX Security Symposium*, San Jose, CA, 1996, pp. 77–90.
- [2] U.S. Department of Defense, "DoD 5220.22-M: National Industrial Security Program Operating Manual," Washington, DC: DoD, 2006.
- [3] M. Wei, L. M. Grupp, F. E. Spada, and S. Swanson, "Reliably Erasing Data from Flash-Based Solid State Drives," in *Proc. USENIX FAST*, San Jose, CA, 2011, pp. 105–117. **[Primary empirical foundation of AAD-50]**
- [4] M. M. Hasan and B. Ray, "Data Recovery from 'Scrubbed' NAND Flash Storage: Need for Analog Sanitization," in *Proc. 29th USENIX Security Symposium (USENIX Security '20)*, 2020. **[Proves data recoverable after digital sanitization via analog threshold voltage analysis — directly justifies AAD-50 multi-cycle Phase B overwrite]**
- [5] C. Abutaleb, J. Bhatt, and R. Panda, "Remanence Effects in NAND Flash under Forensic Microscopy Conditions," *IEEE Trans. Electron Devices*, vol. 68, no. 4, pp. 1820–1831, Apr. 2021.
- [6] ISO/IEC 27040:2015, *Information Technology — Security Techniques — Storage Security*, ISO, Geneva, 2015.
- [7] NVM Express, Inc., *NVM Express Base Specification, Revision 2.0*, NVM Express, Inc., 2021. [Online]. Available: <https://nvmexpress.org/>
- [8] National Institute of Standards and Technology, "NIST SP 800-88 Rev. 2: Guidelines for Media Sanitization," Gaithersburg, MD: NIST, Sep. 2025.
- [9] IEEE, *IEEE Standard for Sanitizing Storage*, IEEE 2883-2022, IEEE, New York, NY, 2022.
- [10] Y. Abeselom, "The Abeselom ASIC-Direct 50 (AAD-50): A Firmware-Enforced, 50-Cycle Sanitization Specification for NVMe Solid-State Storage Media," Version 1.1, June 2026. [Online]. Available: <https://github.com/yonasabeselom/aad50>